



Failure Analysis: Where Does the Model Break?
OASIS-1 Test Set (n=38) + Real 64 mT (n=23)

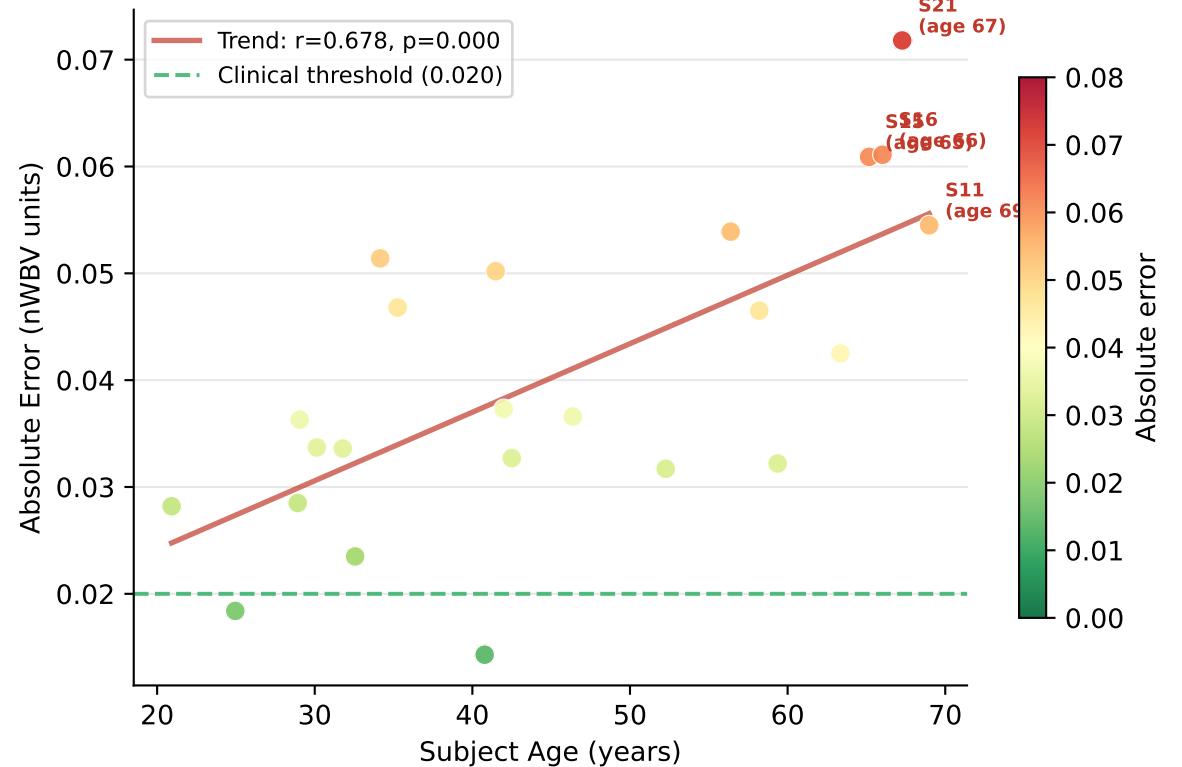
(A) Residual Plot — OASIS-1 Test Set
Model over-predicts for atrophied brains



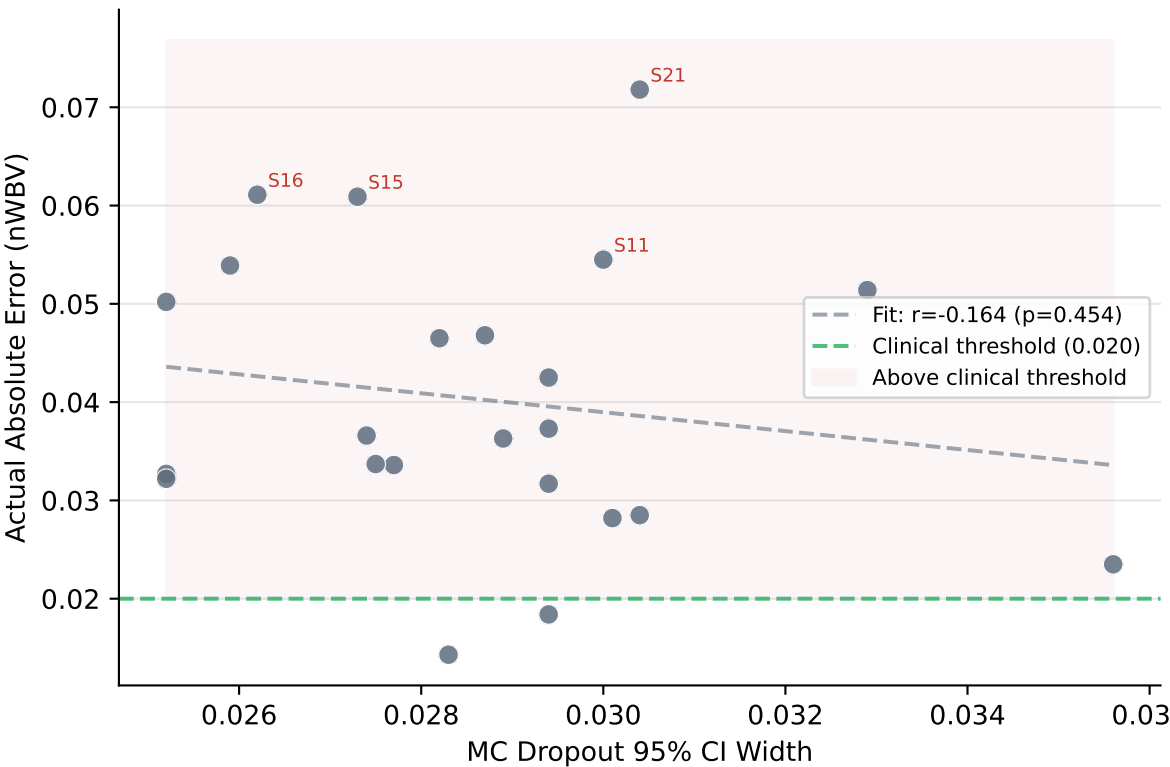
(B) Error by True nWBV Tertile — OASIS
Failure concentrated in atrophied subjects



(C) Error vs Age — Real 64 mT
Older subjects harder ($r=0.678$, $p=0.000$)



(D) CI Width vs Actual Error — Real 64 mT
Poor calibration: $r=-0.164$ ($p=0.454$)



Key failure mode: model predicts near-constant ≈ 0.84 nWBV; fails for atrophied brains (nWBV < 0.76) and older subjects. MC-Dropout CIs do not predict where errors occur (miscalibrated).